

Adaptive Page Allocation: Cost and Replication Analysis

Maksymilian Neumann

May 27, 2025

Abstract

We analyze the performance of the COUNT replication algorithm under varying write probabilities p and threshold values D . We measure two primary metrics: the average cost per request and the average maximum number of replicas. Results are presented via bar charts and heatmaps to illustrate trends and trade-offs.

1 Introduction

In distributed shared memory systems, dynamic replication can reduce read and write costs by adaptively placing copies. The *COUNT* algorithm tracks writes per node and replicates once a threshold D is exceeded. We evaluate how write intensity p and D affect (1) the average cost per request and (2) the maximum number of replicas.

2 Methodology

We simulate a 64-node clique. For each (D, p) pair, we perform 10 000 runs of 65 536 randomly chosen read/write requests (write with probability p , read with probability $1 - p$).

- **Read cost:** 0 if local replica exists, 1 otherwise.
- **Write cost:** $\#replicas - 1$ if requester holds a copy, else $\#replicas$.
- **Replication rule:** replicate to node when its local writes exceed D , incurring D in write cost.

Metrics are averaged across runs.

3 Results

3.1 Average Cost per Request

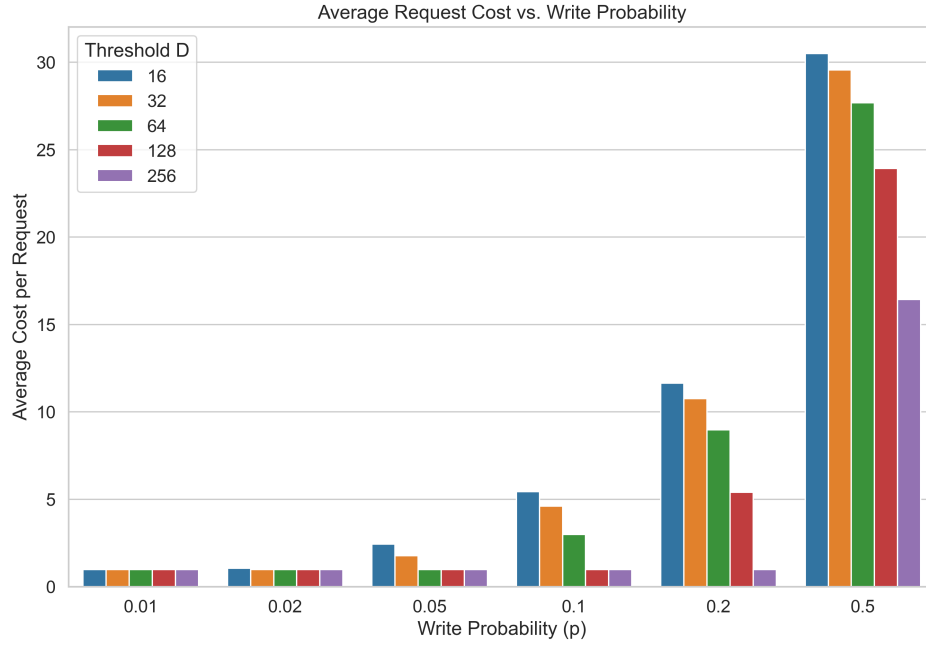


Figure 1: Average request cost vs. write probability p for various thresholds D . Bars show mean cost over 10000 runs.

3.2 Average Maximum Number of Replicas

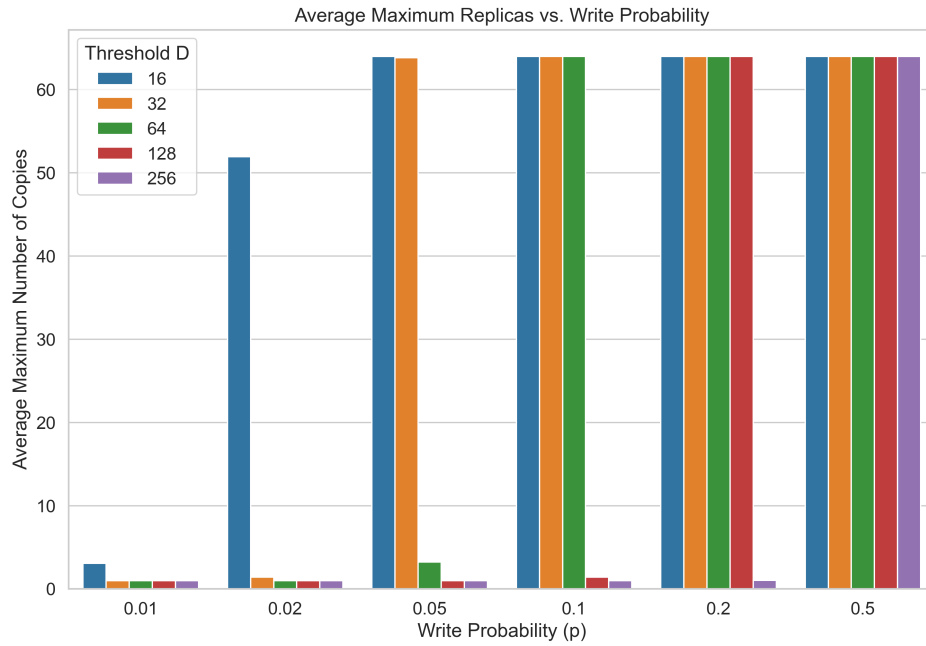


Figure 2: Average maximum number of replicas vs. write probability p for various thresholds D .

3.3 Heatmap of Average Cost



Figure 3: Heatmap showing average cost per request as function of D (rows) and p (columns).

3.4 Heatmap of Average Max Copies

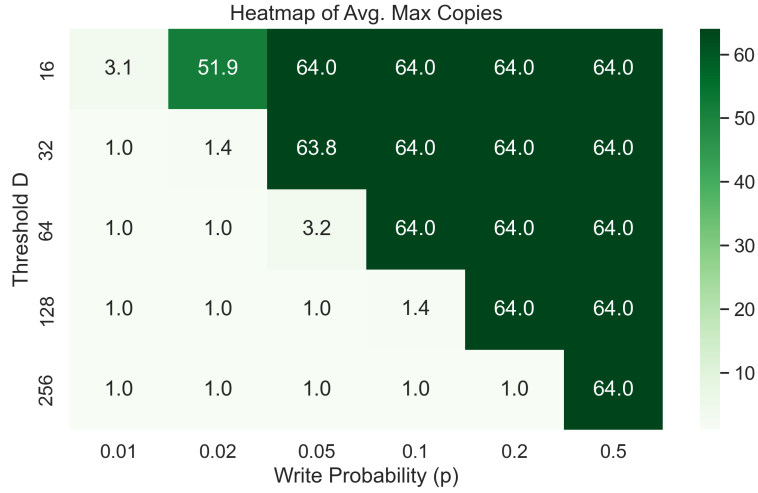


Figure 4: Heatmap showing average maximum replicas as function of D and p .

4 Discussion

- **Cost trends:** Lower D (e.g. 16) replicates aggressively, raising write overhead at high p but reducing read cost at low p . Higher D delays replication, which is beneficial when writes are rare.
- **Replication vs. writes:** For $p \geq 0.1$, most thresholds achieve full replication (64 copies), as shown by the plateau in Fig. 2 and heatmap (4). At moderate p , intermediate D values strike balance.
- **Design implication:** Choose D near the expected write count per node between reads to minimize total cost.